



Deep Research 深度研究

Defending the Crown: Nvidia's Ruthless Strategy to Neutralize Emerging Threats

捍卫王冠：英伟达以冷酷战略化解新兴威胁



Art by Clark Miller.

插画：Clark Miller。



By The Information Staff 作者：The
Information 编辑部

Dec 31, 2025, 9:57am PST

Nvidia is consolidating its dominance in the AI hardware market through a dual strategy of retreating from direct cloud competition and aggressively expanding its footprint in AI inference. By restructuring its DGX Cloud division and striking a \$20 billion licensing and “acqui-hire” deal with startup Groq, the company is shifting from an external service provider to a full-stack architect of the “AI factory.” These moves prioritize protecting its high-margin chip business while neutralizing emerging architectural threats to its market lead.

英伟达正通过一项双重战略巩固其在 AI 硬件市场的主导地位：一方面从直接的云服务竞争中后撤，另一方面积极扩大其在 AI 推理领域的版图。通过重组其 DGX Cloud 业务，并与初创公司 Groq 达成一项价值 200 亿美元的授权与“收购式聘用”交易，该公司正从对外的服务提供商转型为“AI 工厂”的全栈架构师。这些举措优先保护其高利润率的芯片业务，同时化解对其市场领先地位构成的新兴架构性威胁。

We tasked Deep Research to analyze the strategic implications of Nvidia's recent pivot and the Groq acquisition on the AI infrastructure landscape. Here is the response:

我们委托 Deep Research 分析英伟达近期战略转向及收购 Groq 对 AI 基础设施格局的战略影响。以下是其回应：



According to The Information's reporting, Nvidia's recent moves—restructuring its cloud team and striking a massive \$20 billion deal with Groq—are not contradictory but are two sides of the same strategic coin. These actions reveal a company acutely aware of its vulnerabilities and aggressively using its dominant position and vast resources to secure its future against emerging threats. The strategy is to cede ground in high-conflict areas while simultaneously neutralizing threats to its core high-margin business.

根据 The Information 的报道，英伟达近期的一系列动作——重组其云计算团队，并与 Groq 达成一项高达 200 亿美元的巨额交易——并非相互矛盾，而是同一战略硬币的两面。这些举措揭示出，这家公司清醒地认识到自身的脆弱性，并正积极利用其主导性的市场地位和庞大的资源，应对新兴威胁、为未来保驾护航。其策略是在高冲突领域主动让出部分阵地，同时对其核心高利润业务构成的威胁进行阻断和消解。

A Strategic Retreat from Direct Cloud Competition

从正面云竞争中的一次战略性撤退

Our reporting last week revealed that Nvidia is significantly restructuring its DGX Cloud division, backing away from its ambition to compete directly with major cloud providers like Amazon Web Services.¹

我们上周的报道披露，英伟达正在对其 DGX Cloud 业务进行大幅重组，放弃直接与亚马逊云服务（Amazon Web Services）等大型云服务提供商正面竞争的雄心。¹

Initially, Nvidia's DGX Cloud was positioned as a way to offer a premium, optimized environment for AI developers, giving Nvidia a direct relationship with end-users and a new revenue stream. However, the effort faced significant headwinds:

最初，英伟达将 DGX Cloud 打造为一个面向 AI 开发者的高端、深度优化计算环境，使其能够与终端用户建立直接关系，并开辟新的收入来源。然而，这一尝试面临着诸多阻力：

- **Channel Conflict:** The service put Nvidia in direct competition with its largest and most important customers—the cloud providers who buy billions of dollars worth of its GPUs.¹

渠道冲突：这项服务让英伟达与其最大、也是最重要的客户——那些以数十亿美元规模购买其 GPU 的云服务提供商——形成了直接竞争关系。¹

- **Execution Challenges:** Nvidia struggled to attract customers and provide effective technical support, as its service was running on the infrastructure of the very companies it was competing with.¹

执行挑战：英伟达在吸引客户和提供有效的技术支持方面举步维艰，因为其服务运行在它正与之竞争的那些公司所拥有的基础设施之上。¹

- **Shifting Focus:** The cloud team, known as DGX Cloud, is now being merged with Nvidia’s engineering and operations organization. Its primary mission will shift from selling to external customers to serving the massive internal demand from Nvidia’s own engineers for chip development and AI model research.¹

重心转移：被称为 DGX Cloud 的云团队如今正并入英伟达的工程与运营组织。其主要使命将从向外部客户销售，转向满足英伟达自身工程师在芯片开发和 AI 模型研究方面的巨大内部需求。¹

Why it matters: This move is a pragmatic recognition that Nvidia’s strength lies in being the “arms dealer” of the AI revolution, not in fighting trench warfare with established cloud giants. By retreating, Nvidia protects its core chip business and avoids a costly and distracting battle, choosing instead to focus on areas where it has an undisputed technological moat.

为何重要：这一举措务实地承认了英伟达的优势在于充当 AI 革命的“军火商”，而不是与既有的云计算巨头展开堑壕战。通过后撤，英伟达得以保护其核心的芯片业务，避免一场代价高昂且分散注意力的消耗战，转而聚焦于那些它拥有无可争辩技术护城河

Bolstering the AI Inference Moat with Groq

用 Groq 加固 AI 推理护城河

Days after the cloud restructuring news, Nvidia stunned the industry by agreeing to a deal worth approximately \$20 billion to license technology from Groq and hire its CEO Jonathan Ross and other employees.² This was not a standard acquisition but an “acqui-hire” and licensing deal, a structure that may help avoid immediate regulatory scrutiny.³

云服务重组消息公布仅数日后，英伟达就同意达成一项价值约 200 亿美元的交易，从 Groq 获得技术授权，并聘用其首席执行官 Jonathan Ross 及其他员工，震惊了整个行业。² 这并非一桩标准的收购，而是一项“人才并购”（acqui-hire）加技术授权的交易结构，这种做法可能有助于规避监管机构的即时审查。³

Try asking Deep Research to...

不妨试着向“深度研究”提出以下问题.....

I Deep Research



What is next for Apple after the recent executive exodus?

近来多名高管离职之后，苹果接下来将走向何方？



Is Google actually a threat to OpenAI and Nvidia?

谷歌真的会对 OpenAI 和英伟达构成威胁吗？



What major market shifts will define the AI Agent competitive landscape in 2026?

哪些重大的市场变化将定义 2026 年 AI Agent 的竞争格局？



How does Broadcom factor into the AI chip race with Nvidia, AMD and others?

博通在与英伟达、AMD 等公司展开的 AI 芯片竞赛中扮演着怎样的角色？

Groq is a well-funded startup that specializes in chips for AI “inference”—the process of running an already-trained AI model to generate a response. This is a critical area and a potential long-term vulnerability for Nvidia:

Groq 是一家资金雄厚的初创公司，专注于用于 AI“推理”的芯片——这一过程是运行已经训练好的 AI 模型来生成响应。这是一个至关重要的领域，也是英伟达在长期中可能面临的一个薄弱环节：

- **The Rise of Specialized Chips:** While Nvidia’s GPUs are unparalleled for training massive AI models, inference can sometimes be done more efficiently and cheaply on specialized chips. Competitors like Groq, and even Nvidia’s own customers like Google with its TPUs, have been developing such alternatives.^{2,4}

专用芯片的崛起：尽管 Nvidia 的 GPU 在训练大型 AI 模型方面无可匹敌，但在推理阶段，使用专用芯片有时能以更高的效率和更低的成本完成任务。Groq 等竞争对手，甚至包括 Nvidia 自身的客户——例如使用 TPU 的 Google——一直在开发此类替代方案。^{2,4}

- **Groq’s Technology:** Groq’s chip architecture aims to process data faster by keeping it on the processor itself, minimizing the need to move data to and from separate memory chips. This approach could offer significant speed and cost advantages for specific AI applications.³

Groq 的技术：Groq 的芯片架构旨在通过将数据保留在处理器本身上来加快处理速度，从而尽量减少在独立内存芯片之间来回传输数据的需求。这种方法可能为特定 AI 应用带来显著的速度和成本优势。³

- **Neutralizing a Threat:** The deal gives Nvidia access to Groq’s intellectual property and top talent. Huang stated the plan is to integrate Groq’s “low-latency processors into the Nvidia AI factory architecture,” effectively absorbing a key competitor’s technology to strengthen its own offerings for the inference market.²

化解威胁：这笔交易使 Nvidia 获得了对 Groq 知识产权和顶尖人才的访问权。黄仁勋表示，计划将 Groq 的“低延迟处理器整合进 Nvidia 的 AI 工厂架构”，从而在推理市场中有效吸收一家关键竞争对手的技术，强化自身的产品布局。²

Why it matters: This aggressive move shows Nvidia is using its immense cash pile to proactively defend its market share. The company recognizes that as the AI industry matures, the cost of running models (inference) will become as important, if not more so, than the cost of training them. By acquiring Groq’s expertise, Nvidia is betting it can dominate the full spectrum of AI computing, from training to inference, and head off the threat of cheaper, specialized alternatives.

为何重要：这一激进举措表明，英伟达正动用其庞大的现金储备，主动捍卫自身的市场份额。公司已经意识到，随着 AI 行业走向成熟，运行模型的成本（推理）将变得与模型训练成本同等重要，甚至更加重要。通过收购 Groq 的技术专长，英伟达押注于在从训练到推理的全栈 AI 计算领域实现全面主导，从而提前化解来自更便宜、更专业化替代方案的威胁。

The Future: A Paranoid Push for Full-Stack Dominance

未来：偏执式推进的全栈主导之路

Viewed together, these events paint a clear picture of Nvidia’s future strategy, driven by what our reporting describes as CEO Jensen Huang’s productive paranoia.⁵ He is acutely aware that even from a position of unprecedented dominance, the company faces threats from every direction:

将这些事件结合来看，可以清晰地描绘出英伟达的未来战略版图。我们的报道将其归因于首席执行官黄仁勋所具备的“富有成效的偏执”。⁵ 他清醒地意识到，即便公司已处于前所未有的主导地位，仍然在各个方向上面临着威胁：

- **Customer Competition:** Major customers like Google, Amazon, Microsoft, and Meta are all developing their own custom AI chips to reduce their reliance on Nvidia.⁶

客户竞争：谷歌、亚马逊、微软和 Meta 等主要客户都在开发自家的定制 AI 芯片，以减少对英伟达的依赖。⁶

- **Hardware Commoditization:** The rise of viable alternatives for inference workloads could threaten Nvidia's high prices and margins.

硬件商品化：用于推理工作负载的可行替代方案日益增多，可能会威胁英伟达的高定价和高利润率。

- **Ecosystem Control:** Nvidia's ultimate defense is its proprietary CUDA software, which locks developers into its platform. The company is expanding this with more software and services, like Nvidia AI Enterprise, to create a “full stack” that is difficult to leave.⁵

生态系统控制：英伟达的终极防御手段是其专有的 CUDA 软件，它将开发者锁定在其平台上。公司正通过更多软件和服务（如 Nvidia AI Enterprise）来扩展这一优势，打造一个难以脱离的“全栈”体系。⁵

Nvidia's strategy is to become more than just a chip seller. It is building a complete, defensible ecosystem. The cloud retreat shows a disciplined focus on its core mission, while the Groq deal demonstrates a ruthless willingness to spend whatever it takes to acquire technology and talent that can protect its dominance. The future of Nvidia is not just about making the most powerful chips, but about controlling the entire “AI factory,” from hardware and networking to the software that runs on it.

英伟达的战略是成为一家不只是卖芯片的公司。它正在构建一个完整且具备防御性的生态系统。退出云服务领域表明其对核心使命的高度自律，而与 Groq 的交易则展现了其不惜任何代价获取技术和人才、以巩固自身主导地位的冷酷决心。英伟达的未来不仅在于制造最强大的芯片，还在于掌控整个“AI 工厂”——从硬件和网络，到其上运行的软件。

The Information Sources The Information 来源

1. Nvidia Restructures Cloud Team After Retreating From AWS Competition

英伟达在退出与 AWS 的竞争后重组云业务团队

2. Why Nvidia Struck a \$20 Billion Megadeal with Groq

为什么英伟达与 Groq 签订了一项价值 200 亿美元的超级大单

3. Nvidia Agrees to Pay Around \$20 Billion to Access Groq's AI Chip Tech and Hire Its Leaders

英伟达同意支付约 200 亿美元以获得 Groq 的 AI 芯片技术并招募其领导团队

4. Google Convinces OpenAI to Use TPU Chips in Win Against Nvidia

谷歌说服 OpenAI 使用 TPU 芯片，在与英伟达的竞争中取得胜利

5. Nvidia's Jensen Huang Is on Top of the World. So Why Is He Worried

英伟达的黄仁勋正处于巅峰。那么他为何仍然忧心忡忡

6. Why Tech Giants Are Partnering with Broadcom to Break Free from Nvidia

科技巨头为何联手博通以摆脱对英伟达的依赖
